

# **Data Loading, Manipulation, and Visualization**

Introdução Engenharia Informática

---

Mário Antunes

December 01, 2025

Universidade de Aveiro

# Table of Contents i

Data Formats and Structures

Data Loading & Manipulation

Data Visualization

Jupyter Notebooks

Further Reading

# Data Formats and Structures

---

# The Spectrum of Data Structure i

Understanding the underlying organization of data is the first step in the pipeline.

## 1. Structured Data

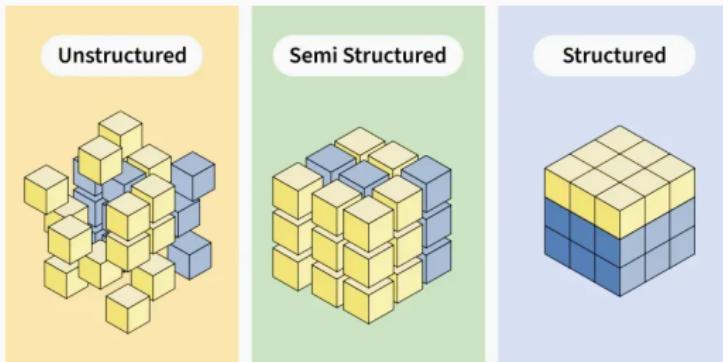
- **Definition:** Data that adheres to a pre-defined data model and is therefore straightforward to analyze.
- **Format:** Rows and columns (Tabular).
- **Schema:** Rigid (Schema-on-write).
- **Examples:** Relational Databases (SQL), Excel files.

## 2. Unstructured Data

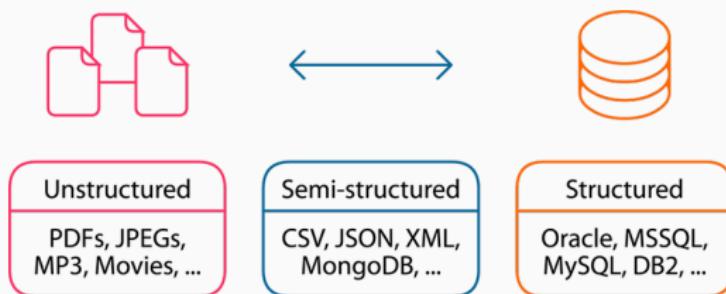
- **Definition:** Information that either does not have a pre-defined data model or is not organized in a pre-defined manner.
- **Format:** Text, Binary, Media.
- **Schema:** None (Schema-on-read).
- **Examples:** PDF documents, Video, Audio, Plain Text emails, Social Media posts.

## 3. Semi-Structured Data

- **Definition:** Data that does not reside in a relational database but has organizational properties that make it easier to analyze. It uses “tags” or “markers” to separate semantic elements and enforce hierarchies.
- **Format:** Hierarchical trees or Key-Value pairs.
- **Examples:** JSON, XML, YAML, NoSQL databases.



**Figure 1:** Illustration of [Un | Semi]Structured organization



**Figure 2:** Examples of [Un | Semi]Structured Files

## CSV (Comma Separated Values)

- **Structure:** Flat text file where lines are rows and commas separate columns.
- **Pros:** Universally supported, extremely lightweight.
- **Cons:** No support for types (everything is a string/number), no nesting/hierarchy.

### XML (eXtensible Markup Language)

- **Structure:** Tree-based structure using custom opening/closing tags.
- **Pros:** Standard for legacy web services (SOAP), supports complex hierarchy and schemas (XSD).
- **Cons:** Verbose (heavy storage footprint due to repeated tags), harder to parse than JSON.

### JSON (JavaScript Object Notation)

- **Structure:** Key-Value pairs using brackets {} for objects and [ ] for arrays.
- **Pros:** The standard for modern Web APIs (REST), native mapping to Python Dictionaries, human-readable.
- **Cons:** Keys are repeated for every record (verbose).

### YAML (YAML Ain't Markup Language)

- **Structure:** Relies on whitespace indentation to define hierarchy.
- **Pros:** The most human-readable format; perfect for configuration files (Docker, Kubernetes).
- **Cons:** Indentation errors can break the file easily; parsing can be slower than JSON.

## BSON (Binary JSON)

- **Structure:** A binary-encoded serialization of JSON-like documents.
- **Pros:** Optimized for speed (traversal) and space; supports types JSON does not (e.g., Date, BinData).
- **Cons:** Not human-readable without a decoder. Used primarily in MongoDB.

## Example: “Employee Record” i

Here is a dataset containing two records represented in all formats. Note how “lists” of employees are handled.

### 1. CSV

No native nesting. “Skills” list requires a custom separator (e.g., pipe |).

```
id,name,skills,active
1,"Jane Doe","Python|SQL",true
2,"Bob Smith","Java|C++",false
```

## Example: “Employee Record” ii

### 2. XML

Requires a root tag to wrap multiple children.

```
<employees>
    <employee id="1">
        <name>Jane Doe</name>
        <skills><skill>Python</skill><skill>SQL</skill>
        </skills>
        <active>true</active>
    </employee>
    <employee id="2">
        <name>Bob Smith</name>
        <skills><skill>Java</skill><skill>C++</skill>
        </skills>
        <active>false</active>
    </employee>
</employees>
```

## Example: “Employee Record” iii

### 3. JSON

JSON uses square brackets [ ] to denote a list.

```
{"employees": [{"  
    "id": 1, "name": "Jane Doe",  
    "skills": ["Python", "SQL"], "active": true  
},  
{  
    "id": 2, "name": "Bob Smith",  
    "skills": ["Java", "C++"], "active": false  
}]  
}
```

## Example: “Employee Record” iv

---

### 4. YAML

YAML uses dashes – to denote list items.

```
employees:
  - id: 1
    name: Jane Doe
    skills: [Python, SQL]
    active: true
  - id: 2
    name: Bob Smith
    skills:
      - Java
      - C++
    active: false
```

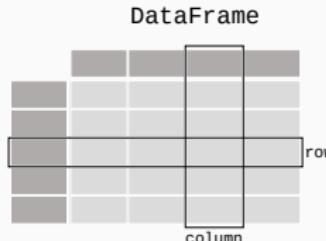
## Data Loading & Manipulation

---

# The DataFrame Concept

A **DataFrame** is the central data structure in data science (used by R, Pandas, Polars, Spark).

- **Conceptual Model:** An in-memory spreadsheet.
- **Structure:**
  - **Index:** Labels for rows (axis 0).
  - **Columns:** Labels for variables (axis 1).
  - **Cells:** The intersection holding the data.
- **Homogeneity:** A single column usually holds a single data type (e.g., all Integers), but different columns can hold different types.



- **Core Philosophy:** Eager execution. Code runs line-by-line immediately, making debugging and exploration intuitive.
- **Key Architecture:**
  - **Single-Threaded:** Primarily runs on a single CPU core.
  - **Index-Based:** Relies heavily on explicit row labels (Indices) for data alignment, which is crucial for time-series analysis.
- **Why use it?**
  - **Unrivaled Ecosystem:** It is the default input for Scikit-Learn, Matplotlib, and thousands of other libraries.
  - **Maturity:** If a data problem exists, there is a StackOverflow answer for how to solve it in Pandas.

- **Core Philosophy:** Lazy Evaluation. It builds an optimized query plan before execution to minimize work and memory usage.
- **Key Architecture:**
  - **Multi-Threaded (Rust):** Written in Rust to bypass Python limitations, utilizing **all available CPU cores** for parallel processing.
  - **Apache Arrow:** Uses a columnar memory format that allows for zero-copy data transfer and efficient caching.
- **Why use it?**
  - **Performance:** Significantly faster than Pandas on large datasets (10x-100x speedups are common).
  - **Streaming:** Can process datasets larger than your computer's RAM.

# Library Comparison: Pandas vs. Polars

Feature	Pandas	Polars
<b>History</b>	Created in 2008. The industry standard.	Newer (2020s). Built for performance.
<b>Backend Execution</b>	Python/C / Cython. <b>Eager:</b> Runs line-by-line immediately.	<b>Rust</b> (Safety & Speed). <b>Lazy &amp; Eager:</b> Can optimize the whole query plan before running.
<b>Parallelism</b>	Single-threaded (mostly).	<b>Multi-threaded</b> (Native parallelization).
<b>Memory</b>	Copies data often (High RAM usage).	<b>Arrow Memory Format</b> (Zero-copy, efficient).

# Data Cleaning: Missing Values (Imputation) i

Data often has gaps (NaN or Null). You must decide to **drop** them or **fill** them. Missing data can be filled in with the mean (continuous values) or mode (discrete values).

- **Pandas Approach:**

```
# Drop rows with any missing values  
df.dropna()
```

```
# Fill with Mean (Imputation)  
mean_val = df['salary'].mean()  
df['salary'].fillna(mean_val, inplace=True)
```

# Data Cleaning: Missing Values (Imputation) ii

- **Polars Approach:**

```
# Drop rows with any missing values
df.drop_nans()

# Fill with Mean
df.with_columns(
    pl.col("salary")
        .fill_null(pl.col("salary")
        .mean()))
)
```

# Data Cleaning: Outliers with IQR i

Outliers are extreme values that deviate from other observations. We use the **Interquartile Range (IQR)** to detect them.

## The Math:

1. **Q1 (25th Percentile):** The middle number between the smallest number and the median.
2. **Q3 (75th Percentile):** The middle number between the median and the highest number.
3. **IQR Formula:**

$$IQR = Q_3 - Q_1$$

## Data Cleaning: Outliers with IQR ii

**The Filter:** Data  $D$  is an outlier if:

$$D < (Q_1 - 1.5 \times IQR) \quad \text{or} \quad D > (Q_3 + 1.5 \times IQR)$$

**Implementation (Pandas):**

```
Q1 = df['age'].quantile(0.25)
Q3 = df['age'].quantile(0.75)
IQR = Q3 - Q1
# Filtering
df_clean = df[~((df['age'] < (Q1 - 1.5 * IQR)) |
(df['age'] > (Q3 + 1.5 * IQR)))]
```

# Data Cleaning: Outliers with IQR iii

---

## Implementation (Polars):

```
Q1 = df["age"].quantile(0.25, interpolation="linear")
Q3 = df["age"].quantile(0.75, interpolation="linear")
IQR = Q3 - Q1
# Filtering
df_clean = df.filter(
    pl.col("age").is_between(q1 - 1.5 * iqr, q3 + 1.5 * iqr)
)
```

# Data Scaling i

Scaling the data can help to balance the impact of all variables on the distance calculation and can help to improve the performance of the algorithm.



**Figure 4:** Scaling Data

## Data Scaling ii

---

### A. Min-Max Scaling (Normalization)

- **Goal:** Squeeze data into a range [0, 1].
- **Formula:**

$$X_{norm} = \frac{X - X_{min}}{X_{max} - X_{min}}$$

- **Usage:** Image processing (pixels 0-255), Neural Networks.

### B. Standard Scaling (Z-Score Standardization)

- **Goal:** Center data around Mean=0 with Std Dev=1.

## Data Scaling iii

---

- **Formula:**

$$Z = \frac{X - \mu}{\sigma}$$

- **Usage:** PCA, Logistic Regression, Clustering.

## Summaries (Central Tendency & Dispersion):

- **Mean:** Average.
- **Median:** Middle value (Robust to outliers).
- **Mode:** Most frequent value.
- **Quartiles:** Distribution checkpoints (25%, 50%, 75%).

## Correlation (Pearson):

- Measures linear relationship between  $-1$  (perfect negative) and  $+1$  (perfect positive).
- $0$  means no linear correlation.

## Code Example:

# Statistics: Correlation & Summaries ii

```
# Pandas/Polars Summary  
# Returns count, mean, std,  
# min, 25%, 50%, 75%, max  
print(df.describe())  
  
# Pandas Correlation Matrix  
print(df.corr())
```

# Data Visualization

---

# Matplotlib vs. Seaborn i

---

## Matplotlib:

- **Role:** The engine.
- **Philosophy:** “Make easy things easy and hard things possible.”
- **Control:** Granular control over every axis, tick, line style, and canvas element.
- **Syntax:** Imperative (Step-by-step construction).

## Seaborn:

- **Role:** The interface.
- **Philosophy:** “Draw attractive and informative statistical graphics.”

## Matplotlib vs. Seaborn ii

---

- **Control:** Automates color mapping, legends, and statistical aggregation.
- **Syntax:** Declarative (Describe *what* you want, not how to draw it).

# Best Plots for Analysis Types i

## 1. Analysis: Distribution (Univariate)

- **Goal:** See the spread and shape of data.
- **Plot:** **Histogram** or **KDE** (Kernel Density Estimate).
- **Example:** `sns.histplot(data=df, x="price", kde=True)`

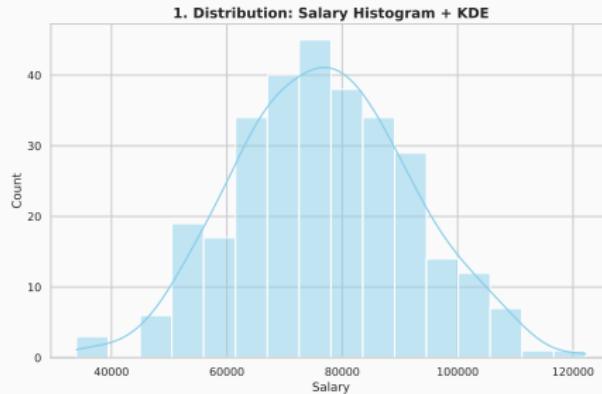


Figure 5: Distribution

# Best Plots for Analysis Types ii

## 2. Analysis: Correlation (Bivariate)

- **Goal:** See how two variables relate.
- **Plot:** Scatter Plot or Heatmap.
- **Example:** `sns.scatterplot(data=df, x="age", y="salary")`



Figure 6: Correlation

# Best Plots for Analysis Types iii

## 3. Analysis: Comparison (Categorical)

- **Goal:** Compare values across groups.
- **Plot:** **Bar Chart** (aggregates) or **Box Plot** (distributions).
- **Example:** `sns.boxplot(data=df, x="department", y="salary")`

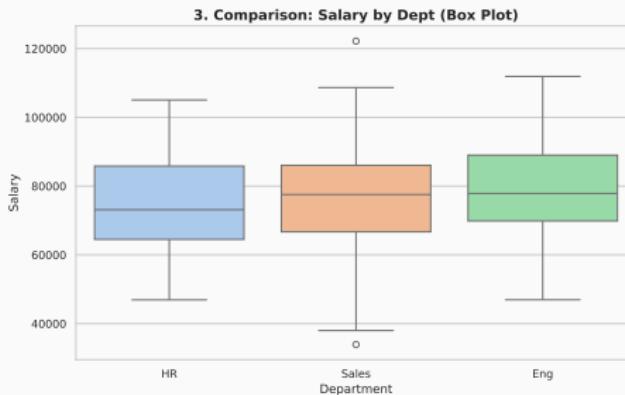


Figure 7: Comparison

# Best Plots for Analysis Types iv

## 4. Analysis: Comparison (Distribution Density)

- **Goal:** Compare the distribution and probability.
- **Plot:** Violin Plot.
- **Example:** `sns.violinplot(data=df,  
x="department", y="salary")`

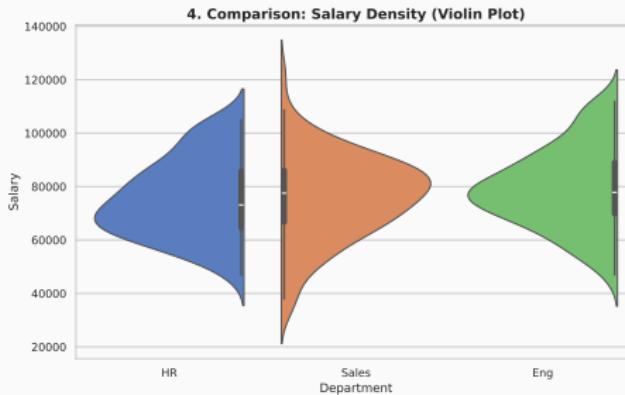


Figure 8: Comparison

# Exporting: Raster vs. Vector

When saving your figures (`plt.savefig`), the extension determines the technology.

Feature	Raster (Bitmap)	Vector
<b>Formats</b>	.png, .jpg, .bmp	.pdf, .svg, .eps
<b>Composition</b>	Grid of colored pixels.	Mathematical formulas (paths, curves).
<b>Scalability</b>	Loses quality when zoomed (pixelates).	Infinite scalability (sharp at any zoom).
<b>File Size</b>	Large for high resolution.	Small (unless containing thousands of scatter points).
<b>Usage</b>	Web, PowerPoints, Quick previews.	<b>Academic Papers, Printing, Posters.</b>

# Jupyter Notebooks

---

## The Interactive Framework i

---

**Jupyter** (JULia, PYThon, R) is a web-based interactive computing platform. It is a JSON document containing an ordered list of input/output cells.

## Key Features

1. **Literate Programming:** Mixes executable code with narrative text (Markdown), equations (“*LaTeX*”), and images.
2. **The Kernel:** The computational engine (e.g., IPython) runs in the background. It maintains the “state” (variables remain in memory between cells).
3. **Visualization:** Plots render inline directly below the code that generated them.

# Effective Usage Workflow i

---

- 1. The “Markdown” Cell** Use for documentation. Supports headers (#), bolding (\*\*), and lists to explain methodology *before* the code.
- 2. The “Code” Cell** Runs Python. The last line of a cell is automatically printed (no `print()` needed).
- 3. Magic Commands**

- `%matplotlib inline`: Embeds plots in the notebook.
- `%timeit`: Runs a line multiple times to benchmark performance.
- `!pip install X`: Runs shell commands to install libraries inside the notebook.

## **Further Reading**

---

## Further Reading

- **Pandas Documentation:** [pandas.pydata.org](https://pandas.pydata.org) - The definitive guide.
- **Polars User Guide:** [pola.rs](https://pola.rs) - Essential for high-performance data manipulation.
- **Seaborn Gallery:** [seaborn.pydata.org](https://seaborn.pydata.org) - Visual inspiration for plots.
- **Matplotlib Anatomy of a Plot:** [matplotlib.org](https://matplotlib.org) - Understanding the object hierarchy.
- **Project Jupyter:** [jupyter.org](https://jupyter.org)